

Final Report: Delivering Elite European Football Analytics

DSCI 591 Capstone: Trilemma Foundation

Ngoc Minh Quan Hoang Teem Kwong Rabindranath Duran Pons Li Pu

2026-06-01

1 Executive Summary

We built a fully automated football analytics platform for the Trilemma Foundation on top of the StatsBomb open event dataset (StatsBomb 2024). The platform ingests raw match, event, and lineup data from 16 competitions (3,464 matches, approximately 12.2 million on-ball events) into Google BigQuery, transforms it through a three-layer dbt pipeline into analytics-ready tables, enriches those tables with two machine learning models, and serves the results through a Looker Studio dashboard and a plain-English chatbot.

Three research questions were addressed. For RQ1, unsupervised clustering (PCA + K-Means, $k = 5$) on 13 per-90 features identified five interpretable outfield archetypes — Low Activity, Defensive Anchor, Creative Playmaker, Pressing Forward, and Creative Winger — across 5,759 player-seasons. Tactical role assignment is driven primarily by position, with league context producing secondary but detectable differences in archetype mix. For RQ2, 489 dual-context players (those with 270 minutes in both club and international football) were scored for cross-context consistency; 32% qualify as Elite performers in both settings, while 36% are single-context specialists. For RQ3, a Django-based chatbot powered by a Groq LLM translates natural language questions into BigQuery SQL and returns formatted answers in under ten seconds, making the full dataset queryable without SQL knowledge.

The delivered product is reproducible, fully documented, and deployable with a single `docker compose up` command.

2 Introduction

2.1 Background

The Trilemma Foundation aims to bring data-driven methods to football analysis. StatsBomb open data (StatsBomb 2024) provides one of the richest publicly available sources of match event data, covering player-level on-ball actions — shots, passes, carries, pressures, dribbles, and more — at sub-second granularity. Despite its depth, the dataset is difficult to work with at scale: event records are deeply nested JSON, competitions and seasons are fragmented across many files, and

extracting meaningful player metrics requires complex aggregation logic that is impractical to write ad hoc.

2.2 The Problem

No systematic, production-grade pipeline existed to transform StatsBomb event data into a unified analytical layer accessible to non-technical users. Analysts who could not write SQL against nested schemas were effectively locked out of the data. Trilemma needed a product that could answer tactical questions — which players fit a pressing role, which players perform differently for their country than their club — without requiring custom queries for every request.

2.3 Refined Scientific Objectives

Over the course of the project, the initial problem was refined into three concrete and measurable research questions:

1. **RQ1 — Player archetypes:** What tactical archetypes exist across European football, and how does role distribution differ by position and competition?
2. **RQ2 — Cross-context consistency:** Do players maintain similar tactical profiles when playing for their club versus their national team, and which players specialize in one context?
3. **RQ3 — Natural language access:** Can a non-technical analyst query the full dataset in plain English and receive accurate, data-grounded answers in near real time?

These objectives were chosen because they are directly actionable for a football analytics organization: archetype labelling supports scouting and squad-building decisions, consistency scoring informs national team selection strategy, and the chatbot removes the SQL barrier for day-to-day analytical work.

The data product addresses all three: the Looker Studio dashboard visualizes RQ1 and RQ2 findings interactively, and the chatbot covers RQ3 while also enabling ad-hoc queries against the full mart layer.

3 Data

3.1 Source

All data in this project comes from the StatsBomb open dataset, accessed via the `statsbomppy` Python library (StatsBomb 2024). StatsBomb collects match event data by manually tracking every on-ball action in a match, assigning it a type, timestamp, location (x, y coordinates on a 120×80 pitch grid), and a rich set of type-specific attributes. The open tier covers 16 competitions spanning men’s senior football across Europe, South America, North America, and international tournaments.

3.2 Dataset Structure

Three tables are ingested from the API (Table 1):

Table 1: Overview of ingested StatsBomb tables.

Table	Grain	Scale
<code>matches</code>	One row per match	3,464 matches
<code>events</code>	One row per on-ball action	~12.2 million rows
<code>lineups</code>	One row per player appearance	165,820 rows

The `events` table is the analytical backbone. Action types are dominated by Pass (~3.4M), Ball Receipt (~3.2M), Carry (~2.6M), and Pressure (~1.1M). There are 88,023 Shot events with non-null StatsBomb xG values (mean xG per shot: 0.107). The dataset covers 10,808 distinct players across 308 teams.

3.3 Data Quality and Limitations

Several quality issues and constraints shaped the project:

- **Missing lineup times.** Substitution `to_time` is missing for 110,917 lineup rows, corresponding to starters who played the full match and bench players who never appeared. Minutes played are inferred as 90 (or 120 for extra time) when the field is null.
- **Missing positions.** `position_name` is absent for 34,179 lineup rows, consistent with bench entries with unconfirmed roles. These rows are excluded from per-position analyses.
- **Uneven competition coverage.** La Liga contributes 868 matches with complete seasonal coverage; the UEFA Champions League has scattered single-match history. Analysts should treat early-season or low-coverage competitions as smaller sub-samples.
- **Open data scope.** The free tier does not include all competitions or all seasons available in the paid StatsBomb API. Coverage is curated by StatsBomb and may change with new releases.

3.4 How Data Shaped Design Decisions

The minimum 270-minute playing time threshold (equivalent to three full matches) was introduced to address the instability of per-90 rates for players with very few appearances. Below this threshold, a single exceptional match can dominate a player-season’s metrics and produce unrepresentative rate statistics. This threshold is applied consistently across all intermediate dbt models, the clustering pipeline, and the consistency scoring pipeline. The choice of 270 minutes was validated against the distribution of player-season minute totals: it retains the majority of regular contributors while removing statistical noise from small samples.

Expected goals (xG) assigns a probability to each shot based on location, angle, and context; **per-90** rates normalise event counts to a full match equivalent so players with different minutes

remain comparable. Both are central to the 13-feature profiles used in clustering and consistency scoring.

The nested JSON structure of the `events` table required flattening at ingestion time. Columns containing dictionary objects are normalized using `pd.json_normalize` before writing to Parquet, and all object-type columns are cast to strings to ensure compatibility with the Apache Arrow schema expected by BigQuery.

4 Data Science Methods

Our analytical work addresses *how* players play, not just output volume. Two models sit on the same 13 per-90 StatsBomb features (attack, progression, defending, efficiency): **PCA + K-Means** for tactical archetypes and **within-context z-score scoring** for club vs national consistency. Full feature definitions, preprocessing steps, and formulas are in [docs/appendix/](#) and [docs/consistency.md](#).

4.1 Player Archetype Clustering

We cluster player-seasons from `int_player_season_stats` after applying a **270-minute** minimum per season ($n = 5,469$ outfield; 5,759 including goalkeepers). **Goalkeepers** are excluded from K-Means and assigned a fixed label because their event profiles are structurally unlike outfield players.

Feature selection. Thirteen per-90 metrics span attacking output (`xg_per_90`, `shots_per_90`), ball progression (`passes_per_90`, `passes_att_third_per_90`, `pass_completion_pct`, `carries_per_90`, `dribbles_per_90`), defensive work (`pressures_per_90`, `interceptions_per_90`, `blocks_per_90`, `clearances_per_90`, `duels_per_90`), and shot efficiency (`xg_per_shot`). The same 13 features are used in clustering, consistency scoring, and the dashboard so labels stay aligned across the product.

Preprocessing and modelling: missing values are median-imputed to preserve the zero-skewed shape of defensive counts; values above the 99th percentile are clipped column-wise to limit outlier influence; features are then standardised. PCA retains components explaining 80% of variance (typically six). **PC1** captures overall involvement volume; **PC2** contrasts defensive actions against attacking output. K-Means with $k = 5$ is applied to the PCA-reduced space. We chose five outfield clusters over slightly tighter statistical fits with fewer groups (Notebook 05) because five labels are easier for coaches to interpret. K-Means was preferred over hierarchical methods for scale and reproducibility on 5,000+ player-seasons.

The resulting archetypes (Table 2) feed the Looker cluster explorer and chatbot marts. Archetypes describe *how a player is used in a given season*, not a fixed career identity — for example, Lionel Messi is mostly classified as Creative Winger but as Pressing Forward during Copa América 2024, consistent with Argentina’s more compact national-team system in that tournament.

Table 2: Player archetype distribution across all player-seasons.

Archetype	Player-seasons	Tactical profile
Low Activity	1,963	Low volume across all metrics; squad players and limited roles
Defensive Anchor	1,470	High clearances, interceptions, duels; low attacking output
Creative Playmaker	807	High pass completion, progressive passes, carries; attacking midfielders
Pressing Forward	766	High pressures, shots, and xG; high-energy forwards and pressing wingers
Creative Winger	463	High dribbles, carries, shots; direct wide attackers
Goalkeeper	290	Assigned by position; excluded from K-Means

4.2 Club vs National Consistency Scoring

For **489 dual-context players** with 270 minutes in both club and national football, we standardise each of the 13 features **within context**: club z-scores compare a player only to other club players; national z-scores only to other national-team players. We then compute PCA-weighted **performance scores** (strength vs peers in each setting), a **consistency score** (1 minus the mean absolute z-gap across features — profile similarity, not quality), and four **performance quadrants** by median splits on club and national performance scores (Table 3).

Table 3: Quadrant assignment logic from `src/ml/consistency.py`.

Club vs median	National vs median	Quadrant
median	median	Elite
median	< median	Club Specialist
< median	median	International Specialist
< median	< median	Underperformer

Progressive passing into the attacking third best separates club-focused from national-focused specialists in exploratory comparisons. Complete formulas and dashboard field definitions: [docs/appendix/consistency-explorer.md](#).

4.3 Data Pipeline

The product runs on a fully automated ELT stack: StatsBomb API → GCS → BigQuery → dbt (staging / intermediate / marts) → Python ML scripts → Django app serving Looker Studio and the Groq chatbot. Dagster orchestrates weekly full runs and hourly refreshes when new data lands in GCS (Figure 1). Pipeline architecture, deployment, and maintenance: [HANDOVER.md](#).




Figure 1: Pipeline architecture: StatsBomb open data flows through GCS and BigQuery, is transformed by dbt into analytics-ready marts, enriched by ML scripts, and served to end-users via Looker Studio and the Django chatbot.

End-to-end validation confirmed aligned row counts (5,759 player-seasons; 489 dual-context players scored). The dbt project follows a three-layer pattern: **staging** cleans and types raw StatsBomb tables; **intermediate** aggregates events to player-match, player-season, and dual-context tables with the 270-minute filter; **marts** join ML outputs for Looker and the chatbot. Materialising marts as BigQuery tables (not views) keeps dashboard and LLM queries fast. Mart schemas: [docs/marts_data_dict.md](#).

5 Data Product and Results

The data product is a unified Django web application combining an embedded Looker Studio dashboard and a natural language chatbot. These forms were chosen because Trilemma’s analysts operate at two levels of inquiry that no single tool handles alone: broad visual exploration and specific ad-hoc questions.

5.1 Results

RQ1 — Role distribution. Chi-square tests on top-five European leagues show **position drives archetype assignment** more strongly than competition (Notebook 07; full heatmaps and statistics in [docs/notebooks.md](#)). Cramér’s V for position × cluster is materially larger than for competition × cluster. In row-normalised heatmaps:

- **Center backs** concentrate in Defensive Anchor and Low Activity; attacking clusters are rare.
- **Forwards and attacking mids** skew toward Pressing Forward and Creative Winger.

- **Central mids** spread across Creative Playmaker, Pressing Forward, and Low Activity, reflecting overlapping box-to-box roles.
- **Full backs** often land in Creative Winger or Creative Playmaker when their profiles emphasise carries and progressive passing.

League effects are secondary but visible (Table 4): Serie A has the largest Low Activity share among the top five; Bundesliga and La Liga show relatively more Pressing Forwards. The same position can shift ± 5 –15 percentage points in archetype share between leagues. Dashboard users should filter by position before comparing leagues.

Table 4: Outfield cluster counts by competition in `mart_player_clusters` (BigQuery). Goalkeepers excluded.

Competition	Largest archetype (player-seasons)	Second largest	Notable pattern
La Liga	Defensive Anchor (304)	Low Activity (167)	Highest volume; anchors lead
Serie A	Low Activity (240)	Defensive Anchor (138)	Most Low Activity share among top five
Ligue 1	Defensive Anchor (175)	Low Activity (148)	Similar anchor vs low-activity mix
Premier League	Defensive Anchor (169)	Low Activity (163)	Near-even top two clusters
Bundesliga	Defensive Anchor (145)	Low Activity (121)	Smallest top-five sample; more Pressing Forward (83)

RQ2 — Club vs national consistency. Of 489 dual-context players, **32% are Elite** (above peer median in both contexts), **32% Underperformer**, and **36% single-context specialists** (18% club-only, 18% national-only; Table 5). About two thirds of the cohort rank above or below peer median in **both** settings rather than in only one. Comparing Club Specialist vs International Specialist players, the largest gaps are in **passes into the attacking third per 90**, **interceptions per 90**, and **carries into the final third per 90** — system and role differences, not raw goal output, drive the split.

A stable profile is not the same as a strong one: many players look similar across contexts while ranking below average in both. Defenders show slightly higher mean consistency than forwards (Notebook 08), suggesting attacking roles shift more between club tactics and national setups. The dashboard’s consistency score and context-shift index measure different things; see [docs/appendix/](#).

Table 5: Observed quadrant counts in `analytics.consistency_scores` (BigQuery).

Quadrant	Players	Share
Elite	157	32.1%
Underperformer	156	31.9%
Club Specialist	88	18.0%
International Specialist	88	18.0%

RQ3 — Natural language access. The chatbot returns BigQuery-backed answers to plain-English questions in under ten seconds. It was evaluated informally against representative analyst questions (player rankings by archetype, cross-context comparisons, competition filters) and consistently applied the same eligibility rules as the dashboard, including the 270-minute threshold. Advanced Mode exposes generated SQL for audit; a read-only SQL guard blocks destructive queries.

5.2 Product design

The **dashboard** provides visual access to pre-computed results — archetype distributions, consistency quadrants, and player rankings — filterable by position, competition, and archetype without writing code. It is appropriate for regular reporting and for communicating findings to non-technical stakeholders within Trilemma’s partner organizations.

The **chatbot** (Groq llama-3.3-70b-versatile) addresses questions that fall outside the dashboard’s fixed charts. Analysts type a question in plain English; the LLM generates BigQuery SQL, executes it against the mart layer, and returns a formatted answer in under ten seconds. An Advanced Mode panel exposes the generated SQL, allowing technical users to audit the LLM’s reasoning and verify that correct filters (e.g., the 270-minute minimum, per-90 normalizations) were applied. A custom `is_safe_sql` guard restricts all LLM output to read-only SELECT statements, preventing prompt injection and unauthorised database writes — an important safeguard when non-technical users interact with production data through natural language.

Strengths: the Dagster-automated pipeline refreshes all data weekly with zero manual intervention, so both interfaces always reflect current StatsBomb releases. The product is fully reproducible from a single `docker compose up` and documented for independent use by a new Trilemma team member.

Limitations: the chatbot handles single-step queries; multi-step reasoning (e.g., “track this player’s archetype across five seasons and then compare against their national team peers”) is not yet supported. Both interfaces are scoped to men’s competitions only — a women’s parallel dbt branch is the most impactful near-term extension. Pre-defined dashboard charts cannot be customized without Looker Studio access; highly specific visual analyses still require a technical user.




Figure 2: Chatbot interface demonstrating a natural language query and Advanced Mode SQL output.




Figure 3: Embedded Looker Studio dashboard providing archetype and consistency explorer views.

6 Conclusions and Recommendations

The original problem was that Trilemma Foundation analysts had no scalable way to extract tactical insights from StatsBomb event data without writing complex SQL. This project addressed that problem by delivering three research products on a single automated pipeline.

RQ1 established that five tactical archetypes emerge from 5,759 player-seasons across 16 competitions, with position being the dominant driver of cluster assignment. Trilemma can now use archetype labels to profile scouting targets and assess squad balance before a transfer window — replacing ad-hoc SQL with a filterable Looker dashboard. **RQ2** found that 36% of dual-context players specialize in one setting (club or national) while 32% perform above peer median in both, providing a data-grounded tool for national team selection decisions. **RQ3** demonstrated that a Groq LLM reliably translates plain English questions into accurate BigQuery SQL, removing the technical barrier for non-technical analysts entirely.

The product addresses the core needs described in the project brief. Remaining limitations are: men’s-only data scope due to feature distribution differences in women’s football, uneven competition coverage in the StatsBomb open tier, and the chatbot’s single-step reasoning constraint. All metrics are relative to peers within each competition or context, so analysts should interpret archetype and quadrant labels as tactical profiles rather than absolute player quality rankings.

Future work should prioritize: a women’s parallel dbt pipeline branch; multi-turn LLM reasoning with session memory for follow-up questions; integration with the paid StatsBomb API for real-time data; and extending the consistency framework to track how individual archetypes evolve across seasons. A key lesson from the project is that product form matters as much as model quality — combining a visual dashboard with a conversational interface gives analysts coverage that neither tool provides independently.

7 Appendix: Supplementary Documentation

Detailed formulas, pipeline architecture, full RQ1/RQ2 analysis notebooks, and mart schemas live in the repository rather than in this report:

Topic	Document
Dashboard metrics and quadrants	docs/appendix/
Consistency pipeline	docs/consistency.md
RQ1/RQ2 notebooks and exports	docs/notebooks.md
ELT architecture and deployment	HANDOVER.md
Mart table schemas	docs/marts_data_dict.md

8 References

StatsBomb. 2024. “StatsBomb Open Data.” <https://github.com/statsbomb/open-data>.